European Management Journal 44 (2026) 550–556

Contents lists available at ScienceDirect

European Management Journal

journal homepage: www.elsevier.com/locate/emj

Insights from the Field

Beyond AI disclosure: Claim accountability and responsible research in scholarly publishing

Ward van Zoonen [a,b,*], Anna Morgan-Thomas [c,d], Aizhan Tursunbayeva [e]

[a] Jyväskylä School of Business and Economics, University of Jyväskylä, Seminaarinkatu 15, Jyväskylän yliopisto 21, Jyväskylä, 40014, Finland
[b] Department Communication Science, Vrije Universiteit Amsterdam, the Netherlands
[c] Management, Adam Smith Business School, University of Glasgow, Glasgow, G128QQ, United Kingdom
[d] Strategy and International Business, Kozminski University, 57/59 ul. Jagiellonska, 03-301, Warsaw, Poland
[e] Dipartimento di Studi Aziendali ed Economici, University of Naples Parthenope, Naples, Italy

ARTICLE INFO

Keywords:
AI
Generative AI
GenAI
Scholarly publishing
Peer review
Claim accountability
Research governance
Responsible research

ABSTRACT

Generative artificial intelligence is now part of everyday scholarly practice. It can help researchers polish text, debug code, translate text, and support deductive coding under human supervision. But it can also produce fluent, competent-looking content that no one is positioned to defend. Journal and publisher policies have largely responded to AI through prohibition, mandatory disclosure, or validity-only permissiveness. Each regulates AI use in the production of manuscripts rather than the defensibility of scholarly claims, leaving the central accountability problem intact. We argue that the right unit of governance is the claim, not the tool. Responsible research with AI requires a named human who can reconstruct and defend each claim that enters the scholarly record. We develop a two-threshold framework, in which the confidentiality threshold protects entrusted material in peer review and editorial deliberation, while the judgment threshold protects human responsibility for the reasoning behind scholarly claims. We then translate the framework into role-based self-assessment for authors, reviewers, editors, and publishers, and identify the institutional anchors that prevent it from decaying into boilerplate. AI is the stress test of scholarly publishing. It has exposed weaknesses in how we evaluate, produce, and protect knowledge, and no gain in efficiency justifies a published claim that no human can defend.

1. Introduction

Generative artificial intelligence (AI) has moved from an object of study to a co-producer of scholarly work (Cornelissen et al., 2024; Yoo, 2024). Under clear human supervision, AI can polish prose, debug code, translate text, and support deductive coding (Dwivedi et al., 2024; Susarla et al., 2023). Used less carefully, the same tools outsource judgment and critical thinking, fabricate information, and substitute automated synthesis for scholarly reasoning (Hannigan et al., 2024; Lindebaum & Fleming, 2024). Either way, researchers now use large language models in ways that current journal policies do not anticipate or detect (Kousha, 2024; Kwon, 2025; Naddaf, 2026). The gap between what journals require, what authors declare, and what enters the published literature has become structural (Goyanes et al., 2025; Yao et al., 2025), and the rapid diffusion of AI is reshaping the scientific workflow (Delios et al., 2025). The governance arrangements of scholarly publishing assume scarcity in knowledge production. They treat the slow, visible stages of drafting and revision as evidence of steady effort (Aguinis et al., 2020; Grimes et al., 2023). AI has challenged that assumption (Susarla et al., 2023), and existing regulatory arrangements have not kept pace with how scholars work (Andrews et al., 2026; Cleland et al., 2026; Stollberger et al., 2025).

Journals have responded with three broad approaches. *Prohibition* treats AI as a categorical risk and excludes it from specified tasks (Gatrell et al., 2024; Nature, 2023; Thorp, 2023). *Disclosure* equates declaration with compliance and asks authors to state where AI was used (Goyanes et al., 2025; Lund et al., 2023). *Validity-only permissiveness* treats methodological soundness as sufficient and declines to police production (Andrews et al., 2026). Each approach governs the process rather than the claims. Each asks whether a tool was used, whether its use was declared, or whether the resulting text looks competent. None asks whether a named human can defend what the machine has shaped.

* Corresponding author. Jyväskylä School of Business and Economics, University of Jyväskylä, Seminaarinkatu 15, Jyväskylän yliopisto 21, Jyväskylä, 40014, Finland.

E-mail address: ward.w.vanzoonen@jyu.fi (W. van Zoonen).

https://doi.org/10.1016/j.emj.2026.06.001
Received 1 May 2026; Received in revised form 1 June 2026; Accepted 7 June 2026
Available online 11 June 2026

All three fail at the analytical level. Accountability requires more than a disclosure form and more specificity than a blanket prohibition (Andrews et al., 2026; Bao & Zeng, 2025). Structural conditions in contemporary publishing, including publication pressure, review overload, and the systematic rewarding of speed, actively undermine honest assessment (Avital, 2024; Bechky & Davis, 2025; Grimes et al., 2023). A regime that cannot distinguish AI used to polish a sentence from AI used to supply an unverified argument abdicates its regulatory function (Cleland et al., 2026; Hannigan et al., 2024). Current AI governance frameworks manage the appearance of compliance while scholarly responsibility erodes (Lindebaum & Fleming, 2024).

We therefore shift the regulatory unit from tool disclosure to claim accountability. The question is whether the responsible human can account for the reasoning, evidence, and judgment that appear in the scholarly record. This anchors accountability in what the machine has materially shaped and what the author must be able to defend (Cleland et al., 2026; Lindebaum & Fleming, 2024). We operationalize this position through two thresholds. The confidentiality threshold protects an institutional trust boundary. Entrusted material has either been exposed to external systems or has not. The judgment threshold protects the human locus of interpretive responsibility. It is graded and allows AI to assist scholars without replacing them. The two thresholds pair a firm safeguard for confidentiality with a graduated standard for human judgment, in a form that survives the obsolescence of any specific tool.

We seek to contribute to the scholarly community's efforts to navigate the looming credibility crisis (Cornelissen, 2026). First, it reframes the governance debate by moving the regulatory unit from process to claim, from compliance to accountability, and from tool governance to claim governance (Andrews et al., 2026; Cleland et al., 2026). Second, it specifies the two thresholds and distinguishes their logics. Confidentiality requires a hard stop. Judgment requires ongoing assessment. The resulting framework remains workable as tools change, where static permissions matrices fail (Andrews et al., 2026; Stollberger et al., 2025). Third, it translates the framework into role-based self-assessment for authors, reviewers, editors, and publishers, and identifies the institutional anchors that keep the instrument from decaying into the ritual it replaces. Role-based alignment ties responsibility to the epistemic obligations of each position (Gatrell et al., 2024; Sarker et al., 2024). The framework gives concrete form to calls for responsible research with AI (Roberson, 2026) and guides what may be considered (in)appropriate use of AI.

2. AI as an amplifier

AI can become either a powerful instrument for advancing science or an anchor that drags scholarly work further into existing pathologies (Cornelissen, 2026). The difference lies less in the tool itself than in the scholarly and institutional conditions under which it is used. Under clear human control, AI may support the craft of research by helping scholars refine prose, debug code, translate text, identify relevant literature, or apply deductively specified coding categories. Used uncritically, the same systems can generate analyses, synthesize literatures, interpret data, and formulate arguments in ways that displace the judgment that makes research credible. Hence, to put this in clear terms: AI may support scholarship when it refines or assists work the scholar controls, but it undermines scholarship when it supplies the analysis, interpretation, or argument that the scholar is expected to provide. Because the presence or absence of a tool is a fact about process and not a certificate of quality, the most consequential question is never directly asked: whether the responsible authors can justify the work (Cleland et al., 2026; Lindebaum & Fleming, 2024).

The governance gap has structural roots that predate AI. What AI amplifies depends on what the institutional environment already rewards (Aguinis et al., 2020; Bechky & Davis, 2025; Susarla et al., 2023). Publication pressure, metric-driven evaluation, review overload, and the prioritization of throughput over interpretive depth shaped the research environment long before large language models became available (Delios et al., 2025; Morgan-Thomas et al., 2024; Ramani et al., 2022). AI did not create these incentives. It lowered the cost of satisfying them. By compressing the time required for drafting and synthesis, the technology makes it easier to generate publication-ready prose from weakly interrogated material, producing output that resembles well-grounded work without the epistemic labor that would secure it (Hannigan et al., 2024; Kulkarni et al., 2024; Nguyen & Welch, 2026; van Quaquebeke et al., 2025). AI-assisted researchers may publish more frequently and attract higher citation counts, while the collective breadth and depth of inquiry contract as attention concentrates on established, data-rich domains (Hao et al., 2026). The result is not simply more research, but a higher risk of sameness, superficial differentiation, and fluent claims whose warrant is difficult to assess (Cornelissen, 2026).

Two regulatory implications follow. First, framings that locate the problem in individual virtue, such as whether researchers have used AI responsibly or declared its use honestly, divert attention from the institutional conditions that shape those choices (McNealis, 2026; Renkema & Tursunbayeva, 2024). Second, accounts that treat AI as a straightforward productivity ally understate the same point. Acceleration is not neutral when the faster production of plausible prose is exactly what weak institutional incentives reward (Grimes et al., 2023; Hao et al., 2026). AI is therefore best understood as an amplifier rather than an origin. It can amplify scholarly craft when it remains subordinate to human judgment, but it can also amplify the production of defensible-looking outputs whose claims no one has adequately checked. This shifts the governance question from whether AI is present to whether AI-shaped claims remain defensible (Andrews et al., 2026; Cleland et al., 2026).

3. Three governance responses to AI

Journal and publisher policies have largely responded to AI in scholarly publishing through three approaches: prohibition, disclosure, and validity-only permissiveness. Each performs a useful function, but each remains focused on the production process rather than on the defensibility of claims. Prohibition treats AI as a source of risk and excludes it from defined tasks, most commonly peer review (Gatrell et al., 2024; Stollberger et al., 2025; Thorp, 2023). The motivation is legitimate. Hallucinated references, fabricated citations, and confidentiality breaches are real threats to the integrity of scholarly work (Ebadi et al., 2025; Wu et al., 2024). The key limitation is misidentification. A manuscript drafted without any technological assistance can misrepresent evidence or delegate interpretation to unread secondary sources. A reviewer who uses AI to draft independent notes can remain fully accountable for the evaluative judgment (Shmueli & Ray, 2024). Therefore, prohibiting the tool does not secure accountability for the claim.

The second response, disclosure, has consolidated rapidly across leading social science journals (Goyanes et al., 2025; Roberson, 2026). Indeed, over 60% of leading social science journals mandate disclosure (Goyanes et al., 2025). Its appeal is that author declarations provide journals with a boundary without requiring them to rebuild their review infrastructure. The logic is visible in formal guidance that asks authors to identify the AI tool used, describe the relevant manuscript content, and affirm responsibility for generated material (see, e.g., COPE, 2023; JAMA Network, 2023). Two problems undercut this advantage. First, the declaration is silent on evidence. The statement records an interaction but tells the reader nothing about whether a quotation was checked, an interpretation is grounded, or a reviewer delegated evaluative judgment (Cleland et al., 2026; Nguyen & Welch, 2026). Second, the model generates perverse incentives. When transparent acknowledgment is read as a signal of diminished merit, the regime rewards strategic silence and plausible retrospective narratives over honest practice (Bao & Zeng, 2025). The empirical pattern is by now legible. The declared use of AI focuses on routine editing (Kousha, 2024). Undeclared use concentrates on introductions and literature reviews, where epistemic grounding is most vital (Yao et al., 2025). Naddaf (2026) notes that more than half of surveyed researchers now use AI to draft peer-review reports, often in violation of explicit guidance. Disclosure, therefore, manages the appearance of compliance while the substance of the problem passes through.

A third response goes further still. Validity-only permissiveness holds that the significance and soundness of a claim are the only criteria that should matter, regardless of production (Andrews et al., 2026). The corrective is useful. It avoids fetishizing the purity of the writing process and refocuses attention on the contribution. Validity, however, is not self-certifying. A fluent argument can be assembled from sources that were never read and data that were never checked. The resulting text may be indistinguishable from warranted scholarship to any reader who has not performed the underlying labor (Davison et al., 2023; Hannigan et al., 2024; Ngwenyama & Rowe, 2024). Plausibility is not justification, and the question of who remains accountable for the conclusions is left unanswered.

Despite their differences, the three responses are administratively expedient, analytically shallow, and share a common defect. Each regulates the production process rather than the defensibility of claims (Andrews et al., 2026; Bao & Zeng, 2025). A tick-box, a disclosure statement, or a comparison against a list of permitted systems makes the presence of tools trivially verifiable at the cost of a form, not the cost of expert judgment (Cleland et al., 2026; Goyanes et al., 2025).

Prohibitions that seek to exclude AI from authorship become untenable as the technology integrates into standard workflows (Andrews et al., 2026; Gatrell et al., 2024; Thorp, 2023). Disclosure regimes prioritize transparency but equate declaration with legitimacy while remaining silent on evidentiary accuracy, interpretive grounding, and authorial responsibility (Bao & Zeng, 2025; Goyanes et al., 2025; Yao et al., 2025). Empirically, disclosure functions as a weak proxy. Declared uses cluster around superficial editing while substantive applications remain undisclosed in epistemically critical sections (Kousha, 2024; Naddaf, 2026; Yao et al., 2025). Validity-focused approaches similarly fall short, either mistaking plausibility for justification or failing to provide operational thresholds to distinguish assistance from substitution (Hannigan et al., 2024; Nguyen & Welch, 2026; Susarla et al., 2023).

What they miss is the failure mode that AI most readily produces: fluent, competent-looking text that no human can reconstruct or defend on challenge (Aguinis et al., 2020; Hao et al., 2026; Van Quaquebeke et al., 2025). Current regimes, therefore, secure procedural legitimacy while leaving claim accountability underdetermined. The question should shift from transparency about production to accountability for the claims themselves.

4. Toward claim accountability

Having identified the common failing of existing regimes, we now develop claim accountability as an alternative. Claim accountability shifts attention from whether AI was used or disclosed to whether the claims that enter the scholarly record can be reconstructed and defended by a named human. This shift changes the unit of evaluation, the test applied to that unit, and the locus of responsibility for the result. The shift from procedural to claim governance is more than relabelling. It changes the unit of evaluation, the test applied to that unit, and the locus of responsibility for the result.

A procedural approach asks whether the prescribed steps were followed and whether the form was filed. A claim regime asks whether the resulting claim can be reconstructed and defended on challenge (Cleland et al., 2026; Lindebaum & Fleming, 2024). Compliance treats the absence of declared violation as evidence of integrity. Accountability treats it as evidence of nothing in particular. Integrity is established by what an author can stand behind (Bao & Zeng, 2025; Hosseini et al., 2023). Similarly, tool governance tracks which software touched the manuscript. Claim governance tests whether each assertion in the record is one a responsible author can own (Andrews et al., 2026; Susarla et al., 2023). Table 1 summarises the three contrasts.

Claim accountability rests on two requirements that are already implicit in scholarly practice but insufficiently specified in current AI governance: defensibility and an identifiable locus of responsibility. Defensibility means that the responsible actor can reconstruct the reasoning, verify the evidence, and justify the interpretation behind a claim. The value of scholarly work lies in this capacity for reconstruction and defense, not in the fluency, speed, or volume with which text is produced (Cornelissen et al., 2024; Ngwenyama & Rowe, 2024).

The locus of responsibility concerns who can be called upon to answer for the claim. Scholarly publishing assigns responsibility to identifiable human actors: authors for manuscripts, reviewers for evaluations, editors for decisions, and publishers for the institutional conditions under which these judgments are made (Susarla et al., 2023). This allocation is what makes accountability enforceable. AI systems cannot bear reputational or professional consequences and therefore cannot occupy this role (Lindebaum & Fleming, 2024).

The two requirements function jointly. Defensibility without an identifiable locus produces text that may be reconstructible but is answerable to no one. A locus without defensibility produces a named actor whose claims cannot survive challenge. Either failure breaks the chain that makes scholarly publishing accountable and credible (Cleland et al., 2026; Susarla et al., 2023). A claim is therefore admissible to the scholarly record only if a named human can reconstruct and defend it (Gregor, 2024).

4.1. Two thresholds for claim accountability

We operationalize claim accountability through two thresholds, one for each failure mode (See Fig. 1). The *confidentiality* threshold concerns whether entrusted material is exposed beyond the boundaries within which it was shared (Cleland et al., 2026; Stollberger et al., 2025). The *judgment* threshold concerns whether the substantive basis for a claim remains attributable to a human actor (Lindebaum & Fleming, 2024; Susarla et al., 2023). Both thresholds hold across roles, disciplines, and technologies, and both mark the point at which assistance becomes substitution. Framed this way, governance tracks the invariants of the system, not the evolution of tools (Andrews et al., 2026).

The *confidentiality* threshold protects entrusted material. Scholarly publishing relies on systems of trust that permit authors to submit unpublished work, allow reviewers to provide candid assessments, and enable editors to deliberate without premature exposure (Hosseini et al., 2023; Susarla et al., 2023). These interactions depend on the assumption that entrusted material remains within the boundaries in which it was shared (Cleland et al., 2026; Gatrell et al., 2024). Recent empirical evidence indicates that this boundary is under sustained pressure. AI use

Table 1
From procedural to claim governance.

Dimension	Procedural governance	Claim Governance
Core question	Were the correct procedures followed? Was AI disclosed?	Can the claim be reconstructed and defended by a named human?
Regulatory unit	Process: tool presence, compliance, output form	Claim: defensibility and accountability.
Evaluation logic	Rule adherence and transparency	Justifiability under scrutiny
Dominant failure mode	Ritualized disclosure; strategic non-disclosure; plausible but weakly grounded output	Confidentiality breach; Claims enter the record that no human can reconstruct or defend
Regulatory instruments	Disclosure policies, prohibitions, output checks	Defensibility thresholds, role-based accountability, verification expectations

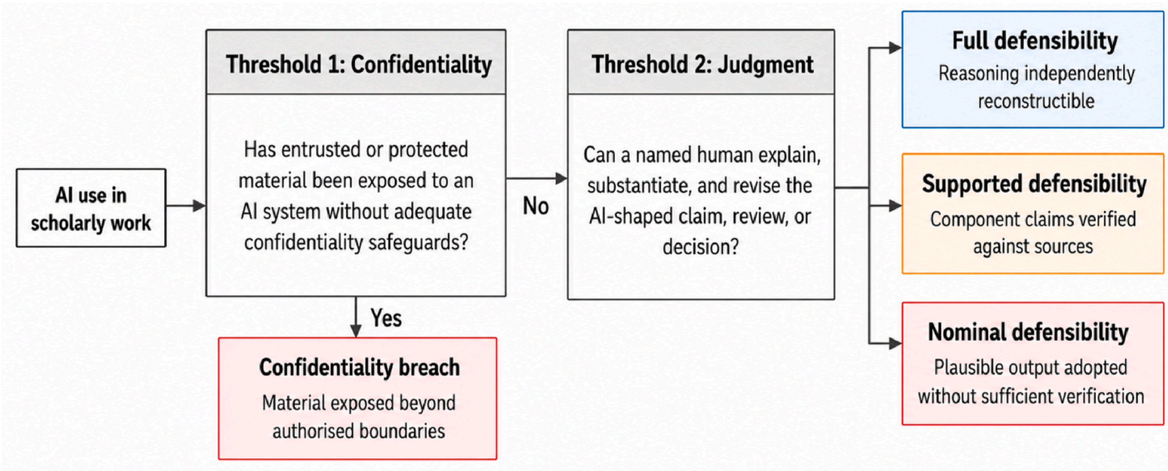

Fig. 1. Confidentiality and judgment: the two-threshold framework.

during peer review occurs frequently despite explicit institutional guidance to the contrary (Naddaf, 2026), and manuscript text may enter external systems during both drafting and reviewing (Ebadi et al., 2025; Wu et al., 2024). The threshold is crossed the moment entrusted material (manuscripts under review, reviewer reports, editorial correspondence) enters a system that cannot guarantee its protection. At that point, the loss of control is definitive, and no further human judgment can repair it (Cleland et al., 2026; Gatrell et al., 2024; Stollberger et al., 2025).

The threshold applies asymmetrically across roles. For reviewers and editors, it is absolute. No portion of an entrusted manuscript, reviewer report, or editorial correspondence should enter an external system, regardless of the actor's estimate of retention risk, because the confidentiality obligation runs to a third party and its breach is not a matter of degree (Ebadi et al., 2025; Naddaf, 2026; Wu et al., 2024). For authors working on their own unpublished manuscripts, no equivalent third-party obligation exists prior to submission. The operative constraint is originality and control. The author must be able to represent the work as genuinely their own and attest that it has not been disseminated in ways that breach the journal's submission conditions (Gatrell et al., 2024). This does not prohibit author-AI interaction during drafting. It requires authors to use systems whose data-handling terms are consistent with those representations, and places the obligation on publishers to certify or provide such infrastructure, rather than leaving individual authors to assess it on a case-by-case basis (Andrews et al., 2026; Naddaf, 2026).

The *judgment* threshold addresses the exercise of judgment and operates through a different logic. Confidentiality is breached by a discrete act of exposure. Judgment is breached by omission when the responsible actor fails to perform the evaluative work inherent to a scholarly role (Cleland et al., 2026; Lindebaum & Fleming, 2024). The question concerns authorship in the substantive sense: who provides the reasons for claims?

The judgment threshold admits AI assistance. Tasks such as literature search, language editing, translation, summarisation, and code generation can support scholarly work without displacing human judgment (Dwivedi et al., 2024; Susarla et al., 2023). Uncritical reliance on AI, however, crosses the judgment threshold. Any use of AI that leaves scholars unable to verify, revise, or defend the claims they advance is incompatible with responsible research. The threshold is crossed when the system supplies what the human is expected to supply. The justification for an argument, the evaluative basis of a review, and the reasoning behind an editorial decision all remain a human responsibility (Cleland et al., 2026; Lindebaum & Fleming, 2024; Shmueli & Ray, 2024). This distinction maps closely onto recent editorial guidance. *Administrative Science Quarterly*, for example, distinguishes AI used to refine prose or debug analytical code from AI used to generate analyses, interpret qualitative data, or develop entire arguments. The former supports scholarship. The latter displaces the judgment that makes research credible.

4.2. Governance framework: hard constraints and soft accountability

The judgment threshold is harder to assess than the confidentiality threshold. The confidentiality threshold is binary. Entrusted material has either been exposed to external systems or has remained secure. No partial exposure falls within acceptable bounds, and no remediation can restore protection once compromised (Cleland et al., 2026; Gatrell et al., 2024; Stollberger et al., 2025). The judgment threshold, by contrast, admits of degrees. It allows AI to inform a decision without necessarily determining its outcome (Drori & Te'eni, 2024; Shmueli & Ray, 2024; Susarla et al., 2023). The primary question is whether the human responsible for a conclusion remains positioned to own, defend, and revise it under scrutiny (Cleland et al., 2026; Lindebaum & Fleming, 2024). Assessment, therefore, shifts from the text or the tool to the actor: from whether AI contributed to whether the responsible actor can explain, defend, and revise the reasoning when challenged. The judgment threshold cannot be enforced by inspecting outputs. It can only be enforced by holding the responsible actor to a defensibility standard at the level of the claim (Shmueli & Ray, 2024; Susarla et al., 2023).

Both mechanisms are necessary. Governance regimes that rely only on hard constraints (such as prohibitions and permissions matrices) can protect confidentiality but collapse when faced with graduated forms of AI assistance (Andrews et al., 2026; Stollberger et al., 2025). Regimes that rely only on soft accountability (disclosure, calls for reflexivity) can accommodate assistance but fail to protect the underlying trust infrastructure (Bao & Zeng, 2025; Goyanes et al., 2025). Durable governance requires both. The two thresholds shift governance from procedural legitimacy to claim accountability by protecting the conditions under which claims remain defensible, and responsibility remains human.

5. Operationalizing claim accountability across roles

Claim accountability manifests differently across academic roles and phases of the publication process. The self-assessment instrument in Table 2 translates the proposed two-threshold framework into role-specific questions that scholars apply at the moment of action rather than at the moment of disclosure. Unlike a compliance or disclosure checklist, the instrument prompts scholars to pause where AI assistance is most likely to shade into substitution (Shmueli & Ray, 2024; Stollberger et al., 2025). It asks whether confidentiality has been

Table 2
Role-based self-assessment instrument.

Role	Confidentiality threshold	Judgement threshold	Academic obligation
Author	Has any portion of the manuscript (including prior drafts or supplementary material) been entered into a system that may retain or reuse the content?	For every interpretive claim, citation, quotation, and inferential step, could I reconstruct the reasoning and defend it from my own reading if the AI output were unavailable?	State defensibility, not declaration. Retain working notes and verification trails that allow a future challenge to be met.
Reviewer	Has any part of the manuscript, my draft reasoning about it, or the review correspondence been entered into an (external) AI system?	Would I stand behind every evaluative point in this report if the AI-generated text disappeared? Are the reasons my own, independently formed from reading the manuscript?	Decline the review rather than breach the confidentiality threshold. Use only publisher-secure infrastructures for any AI-assisted support.
Editor	Have all manuscript-related materials (including reviewer reports and editorial correspondence) remained inside the journal's secure editorial systems at every stage?	Can I state the rationale for this decision in wholly human terms (my reading of the reviews, my reading of the paper, my judgment of fit) without relaying machine-generated reasons?	Document decisional reasons in the journal system. Treat AI recommendations as input, never as grounds.
Publisher	Does the infrastructure we provide to reviewers and editors remove the need to reach for external tools to do their work?	Do our incentive structures credit defensibility (careful reading, verified citation, reasoned disagreement) or do they reward throughput and speed?	Invest in confidentiality-preserving infrastructure. Build post-publication review and audit mechanisms. Reward the interpretive labor that defensibility requires.

preserved (Cleland et al., 2026; Ebadi et al., 2025) and whether the substantive judgment behind a claim, recommendation, or decision remains genuinely theirs (Lindebaum & Fleming, 2024; Susarla et al., 2023). The instrument keeps the unit of analysis anchored to human responsibility (Cleland et al., 2026; Gregor, 2024).

Authors carry primary responsibility for the defensibility of published claims, and this responsibility cannot be transferred to the AI systems that helped produce the text (COPE, 2023; Lindebaum & Fleming, 2024). The guiding question is whether every interpretive claim, citation, quotation, and inferential step in the manuscript could be defended if the author had to reconstruct it from their own reading (Gatrell et al., 2024; Lindebaum & Fleming, 2024). This reframes the author's task from declaration to defensibility. An author who meets this standard has used the technology as an instrument. An author who cannot has substituted it for scholarship (Davison et al., 2023).

Reviewers sit where both thresholds apply with full force. Confidentiality is absolute. The reviewer holds entrusted material that has not yet entered the public record, and no portion of that material should be exposed to an external system, whatever the reviewer's estimate of retention risk (Ebadi et al., 2025; Naddaf, 2026). The judgment threshold is especially stringent for reviewers because their contribution involves evaluative reasoning. A reviewer who would not stand behind every point in their review if the AI output were to disappear has crossed the threshold, even if the review reads competently (Drori & Te'eni, 2024; Sarker et al., 2024).

Editors face the judgment threshold in a form specific to curation. AI systems will plausibly become routine for administrative tasks, completeness checks, reviewer matching, and metadata extraction (Drori & Te'eni, 2024; Shmueli & Ray, 2024). The threshold is crossed when AI supplies the grounds for an editorial decision rather than supports the editor in reaching one. An editor who can articulate the reasons for a recommendation in wholly human terms has preserved the locus. An editor who could not reconstruct the rationale without the tool has ceded it (Kankanhalli, 2024). Editorial engagement with AI remains legitimate as long as the assistance supports the editor in reaching an independent conclusion and does not supply the primary grounds for acceptance, revision, or rejection (Cleland et al., 2026; Gatrell et al., 2024).

Publishers bear responsibility for the incentive environment in which the other three roles work. A self-assessment instrument is not self-enforcing. A journal that publicly commits to accountability while privately rewarding speed creates the conditions under which the judgment threshold will be crossed (Bechky & Davis, 2025; Delios et al., 2025; Kulkarni et al., 2024). The publisher's obligation is to build the anchors that allow individual actors to discharge their duties without becoming martyrs to them: secure manuscript infrastructure, selective probes of AI use in peer review, post-publication review mechanisms, and reward systems that credit defensibility rather than throughput (Gregor, 2024; Yoo, 2024).

6. Discussion

The framework developed here positions AI governance in scholarly publishing as a problem of claim accountability rather than tool management. Its purpose is to preserve the conditions under which scholarly claims remain answerable to human judgment. Existing governance regimes intervene around the claim rather than at it. Prohibition restricts tool use. Disclosure records tool use. Validity-only permissiveness evaluates the apparent quality of the output. Each performs a useful function. None is sufficient where fluent text can mask the absence of accountable reasoning. The proposed two thresholds translate this problem into operational terms. Confidentiality sets a hard boundary around entrusted material. Judgment asks whether a human actor still owns the reasoning AI has helped shape. The shift is deliberately conservative. It introduces no special ethics for AI. It re-specifies existing expectations of authorship, review, and editorial responsibility under changed technological conditions.

The first contribution is to relocate accountability to the claim entered into the published record, the level at which scholarly harm occurs. Disclosure mandates, permissions matrices, and calls for reflexivity all address important aspects of AI use. They do not directly test whether the claim itself remains defensible (Andrews et al., 2026; Goyanes et al., 2025). This matters because disclosure can be observed closely even when the underlying reasoning is weak, unverifiable, or delegated (Lindebaum & Fleming, 2024). Disclosure can also create incentives for strategic silence, in which honest disclosure is treated as a signal of diminished scholarly merit (Bao & Zeng, 2025). Claim-centred governance, therefore, addresses a failure mode that process-centred governance leaves exposed.

The second contribution is to distinguish two forms of accountability failure that require different regulatory logics. Confidentiality concerns the exposure of entrusted material that cannot be partially undone. Judgment is graded because AI may inform scholarly reasoning without replacing it. Critiques of permission matrices show why static rules struggle to govern rapidly changing tools (Andrews et al., 2026; Stollberger et al., 2025). Principle-based accounts of accountability identify what governance should protect without specifying how to enact it (Lindebaum & Fleming, 2024). The two-threshold framework links each threshold to a specific failure mode: a confidentiality breach and a loss of human accountability for claims.

The third contribution is operational. The framework translates claim accountability into role-specific questions because authors, reviewers, editors, and publishers do not face the same AI-related risks or hold the same responsibilities. The relevant question is what kind of responsibility the actor was discharging when AI was used (Gatrell et al., 2024; Sarker et al., 2024). This role-based design moves governance closer to the moment of action, where assistance can become substitution. It also clarifies why institutional support is not an implementation detail. Without a secure infrastructure, credible review mechanisms, and reward systems that prioritize defensibility over throughput, self-assessment will decay into the same ritualized compliance that disclosure regimes already risk producing. AI's effects on scholarly publishing remain anchored in the incentive structures, editorial designs, and collective choices that define what a research community values (Aguinis et al., 2020; Ramani et al., 2022; Shmueli & Ray, 2024). The problem AI amplifies is fundamental, and we all seem to fall for it. But the craft of science was never meant to be a popularity contest won by those best able to dress weak ideas in loud prose. To treat science as a race for polished theory and outputs is to misunderstand our craft. Scholarship advances through the slow and painstaking dialogue between problem-finding and problem-solving, where curiosity is disciplined, and evidence is thoroughly scrutinized.

"Chaos in the Brickyard" makes clear that the problem predates AI (Forscher, 1963). Already in 1963, he warned that science could become flooded with "bricks," stored in ever more "journals," while scholars lost sight of the difference between a pile of outputs and a genuine intellectual edifice. Generative AI renews this danger at scale. It lowers the cost of producing plausible scholarly bricks, but it does not by itself build explanations. That task still requires theory, judgment, evidence, and accountable human authorship.

That said, we realize that our propositions have limitations too. We outline three such limits. The first concerns computational research. In machine-learning papers, bibliometric studies, and parts of computational social science, the machine is not merely an aid to writing or evaluation but part of the method itself. The judgment threshold does not disappear in such cases. It applies at the level of problem framing, model choice, validation, interpretation of outputs, and the claims drawn from them, rather than at every computational step (Hannigan et al., 2024).

The second limitation follows from the graded nature of judgment. Its flexibility is necessary because AI can inform scholarly work without automatically displacing human responsibility, but it also places weight on honest self-assessment and scholars' commitment to intellectual craftsmanship, which disclosure regimes can elicit unreliably (Bao & Zeng, 2025; Lindebaum & Fleming, 2024). The self-assessment instrument is prospective. It tests a capacity at the moment of action by asking whether the actor can defend the claim now, before it enters the record. A researcher who discovers they cannot reconstruct an argument without the AI output and honestly admits as much receives actionable information at the point when correction remains possible. The instrument does not eliminate the dependence on good faith. It relocates the test to the point at which the actor can still recognize failure, rather than relying on external detection after the fact.

The third limit is ritualization. The framework could itself become a boilerplate attestation unless publishers create the conditions under which defensibility can be checked, challenged, and rewarded (Goyanes et al., 2025; Gregor, 2024; Yao et al., 2025; Yoo, 2024). Without those institutional anchors, the framework describes good practice without making it expectable.

7. Conclusion

The central question the scholarly community now faces is whether the human responsible for each claim retains the capacity to account for the reasoning and evidence supporting it (Cleland et al., 2026; Susarla et al., 2023). What makes this question urgent is the possibility that fluent, competent-looking text now enters the scholarly record without any named human in a position to defend it. Disclosure alone cannot detect this failure. The failure lies not in the use of an undeclared tool but in the loss of the human locus of interpretive responsibility. The two-threshold framework proposed protects what remains portable across tool generations: the confidentiality of entrusted material and the accountability of human judgment. AI is the stress test of our scholarly system. It exposes weaknesses in how we already evaluate, produce, and protect knowledge. The integrity of published scientific work will depend less on the tools we permit than on whether we still ask those whose names appear on the page to defend the claims they make.

CRediT authorship contribution statement

Ward van Zoonen: Writing – review & editing, Writing – original draft, Conceptualization. **Anna Morgan-Thomas:** Writing – review & editing, Writing – original draft, Visualization, Conceptualization. **Aizhan Tursunbayeva:** Writing – review & editing, Writing – original draft, Conceptualization.

Declaration of competing interests

The author declares no competing interests.

References

Aguinis, H., Cummings, C., Ramani, R., & Cummings, T. (2020). Reviewing and theorizing the unintended consequences of performance management systems. *Academy of Management Perspectives, 34*(1), 135–154.

Andrews, D. S., Fainshmidt, S., & Gaur, A. (2026). Progress or perish: International business and the adoption of artificial intelligence. *Journal of World Business, 61*, Article 101689.

Avital, M. (2024). Digital transformation of academic publishing: A call for the decentralization and democratization of academic journals. *Journal of the Association for Information Systems, 25*, 172–181.

Bao, A., & Zeng, Y. (2025). AI disclosure, moral shame, and the punishment of honesty. *Accountability in Research*. Advance online publication.

Bechky, B. A., & Davis, G. F. (2025). Resisting the algorithmic management of science: Craft and community after generative AI. *Administrative Science Quarterly, 70*(1), 1–22.

Cleland, J., Driessen, E., Masters, K., Lingard, L., & Maggio, L. A. (2026). When and how to disclose AI use in academic publishing: AMEE guide no. 192. *Medical Teacher, 48*(4), 542–553.

COPE. (2023). *Authorship and AI tools (COPE position statement)*. Committee on Publication Ethics. https://publicationethics.org/cope-position-statements/ai-author.

Cornelissen, J. (2026). The AI con: The commodification of theory and the new and improved credibility crisis. *Organization Studies*. Online-First. 01708406261450467.

Cornelissen, J., Höllerer, M. A., Boxenbaum, E., Faraj, S., & Gehman, J. (2024). Large language models and the future of organization theory. *Organization Theory, 5*, 1–15.

Davison, R. M., Laumer, S., Tarafdar, M., & Wong, L. H. M. (2023). Pickled eggs: Generative AI as research assistant or co-author? *Information Systems Journal, 33*, 989–994.

Delios, A., Tung, R. L., & van Witteloostuijn, A. (2025). How to intelligently embrace generative AI: The first guardrails for the use of GenAI in IB research. *Journal of International Business Studies, 56*, 451–460.

Drori, I., & Te'eni, D. (2024). Human-in-the-loop AI reviewing: Feasibility, opportunities, and risks. *Journal of the Association for Information Systems, 25*, 98–109.

Dwivedi, Y. K., Malik, T., Hughes, L., & Albashrawi, M. A. (2024). Scholarly discourse on GenAI's impact on academic publishing. *Journal of Computer Information Systems*. Advance online publication.

Ebadi, S., Nejadghanbar, H., Salman, A. R., & Khosravi, H. (2025). Exploring the impact of generative AI on peer review: Insights from journal reviewers. *Journal of Academic Ethics, 23*, 1383–1397.

Forscher, B. K. (1963). Chaos in the brickyard. *Science, 142*(3590), 339-339.

Gatrell, C., Muzio, D., Post, C., & Wickert, C. (2024). Here, there and everywhere: On the responsible use of artificial intelligence (AI) in management research and the peer-review process. *Journal of Management Studies, 61*, 739–751.

Goyanes, M., Lopezosa, C., & Piñeiro-Naval, V. (2025). The use of artificial intelligence (AI) in research: A review of author guidelines in leading journals across eight social science disciplines. *Scientometrics, 130*, 3725–3741.

Gregor, S. (2024). Responsible artificial intelligence and journal publishing. *Journal of the Association for Information Systems, 25*, 48–60.

Grimes, M., von Krogh, G., Feuerriegel, S., Rink, F., & Gruber, M. (2023). From scarcity to abundance: Scholars and scholarship in an age of generative artificial intelligence. *Academy of Management Journal, 66*, 1617–1624.

Hannigan, T. R., McCarthy, I. P., & Spicer, A. (2024). Beware of botshit: How to manage the epistemic risks of generative chatbots. *Business Horizons, 67*, 471–486.
Hao, Q., Xu, F., Li, Y., & Evans, J. (2026). Artificial intelligence tools expand scientists' impact but contract science's focus. *Nature, 649*, 1237–1243.
Hosseini, M., Resnik, D. B., & Holmes, K. (2023). The ethics of disclosing the use of artificial intelligence tools in writing scholarly manuscripts. *Research Ethics, 19*(4), 449–465.
JAMA Network. (2023). Use of AI in research and publishing. *JAMA Network*. https://jamanetwork.com/journals/jama/pages/instructions-for-authors#SecUseofAIinPublicationandResearch.
Kankanhalli, A. (2024). Peer review in the age of generative AI. *Journal of the Association for Information Systems, 25*, 76–84.
Kousha, K. (2024). How is ChatGPT acknowledged in academic publications? *Scientometrics, 129*, 7959–7969.
Kulkarni, M., Mantere, S., Vaara, E., van den Broek, E., Pachidi, S., Glaser, V. L., Gehman, J., Petriglieri, G., Lindebaum, D., Cameron, L. D., Rahman, H. A., Islam, G., & Greenwood, M. (2024). The future of research in an artificial intelligence-driven world. *Journal of Management Inquiry, 33*, 207–229.
Kwon, D. (2025). Is it OK for AI to write science papers? Nature survey shows researchers are split. *Nature, 641*, 574–578.
Lindebaum, D., & Fleming, P. (2024). ChatGPT undermines human reflexivity, scientific responsibility and responsible management research. *British Journal of Management, 35*, 566–575.
Lund, B. D., Wang, T., Mannuru, N. R., Nie, B., Shimray, S., & Wang, Z. (2023). ChatGPT and a new academic reality: Artificial intelligence-written research papers and the ethics of the large language models in scholarly publishing. *Journal of the Association for Information Science and Technology, 74*, 570–581.
McNealis, R. (2026). Shame in the machine: Affective accountability and the ethics of AI. *AI & Society, 41*, 403–413.
Morgan-Thomas, A., Tsoukas, S., Dudau, A., & Gąska, P. (2024). Beyond declarations: Metrics, rankings and responsible assessment. *Research Policy, 53*, Article 105093.
Naddaf, M. (2026). More than half of researchers now use AI for peer review—often against guidance. *Nature, 649*, 273–274.
Nature. (2023). Tools such as ChatGPT threaten transparent science; here are our ground rules for their use [Editorial] *Nature* https://www.nature.com/articles/d41586-023-00191-1.
Nguyen, D. C., & Welch, C. (2026). Generative artificial intelligence in qualitative data analysis: Analyzing—or just chatting? *Organizational Research Methods, 29*, 3–39.
Ngwenyama, O., & Rowe, F. (2024). Should we collaborate with AI to conduct literature reviews? Changing epistemic values in a flattening world. *Journal of the Association for Information Systems, 25*, 122–136.
Ramani, R., Aguinis, H., & Coyle-Shapiro, J. (2022). Defining, measuring, and rewarding scholarly impact: Mind the level of analysis. *The Academy of Management Learning and Education, 21*, 470–486.
Renkema, M., & Tursunbayeva, A. (2024). The future of work of academics in the age of artificial intelligence: State-of-the-art and a research roadmap. *Futures, 163*, Article 103453.
Roberson, Q. (2026). Artificial intelligence and responsible research at AMJ. *Academy of Management Journal, 69*(2), 207–211.
Sarker, S., Susarla, A., Gopal, R., & Thatcher, J. B. (2024). Democratizing knowledge creation through human-AI collabouration in academic peer review. *Journal of the Association for Information Systems, 25*(1), 158–171.
Shmueli, G., & Ray, S. (2024). Reimagining the journal editorial process: An AI-augmented versus an AI-driven future. *Journal of the Association for Information Systems, 25*, 60–75.
Stollberger, J., Anand, S., & Dick, P. (2025). Capturing a moving target: Developing research on and with AI for *Human Relations*. *Human Relations, 78*(5), 499–516.
Susarla, A., Gopal, R., Thatcher, J. B., & Sarker, S. (2023). The Janus effect of generative AI: Charting the path for responsible conduct of scholarly activities in information systems. *Information Systems Research, 34*(2), 399–408.
Thorp, H. H. (2023). ChatGPT is fun, but not an author. *Science, 379*(6630), 313.
Van Quaquebeke, N., Tonidandel, S., & Banks, G. C. (2025). Beyond efficiency: How artificial intelligence (AI) will reshape scientific inquiry and the publication process. *The Leadership Quarterly, 36*, Article 101895.
Wu, H., Li, W., Chen, X., & Li, C. (2024). Not just disclosure of generative artificial intelligence like ChatGPT in scientific writing: Peer-review process also needs attention. *International Journal of Surgery, 110*, 5845–5846.
Yao, M., Wei, Y., & Liu, H. (2025). *AI practices and ethical concerns: An analysis of undeclared uses of AI in published research articles. Ethics & Behavior*. Advance online publication.
Yoo, Y. (2024). Evolving epistemic infrastructure: The role of scientific journals in the age of generative AI. *Journal of the Association for Information Systems, 25*, 137–144.